

Project: EXPLORATORY DATA ANALYSIS

INTRODUCTION

This project is to perform an exploratory data analysis for the **10 academy week 0** challenge. The analysis will utilize by using [Solar radiation measurement data](#). The dataset contains the values for solar radiation, air temperature, relative humidity, barometric pressure, precipitation, wind speed, and wind direction, cleaned and soiled radiance sensor (soiling measurement) and cleaning events.

I will use **python libraries** to do the the exploratory data analysis. Firstly, I gather the data hence the data provided by 10 Academy is already downloaded. The data folder contains three CSV files, each representing solar radiation data from an African country. Here is the list of data:

- **benin-malanville.csv**
- **sierraleone-bumbuna.csv**
- **togo-dapaong_qc.csv**

Then, I will convert each country's data into a separate **Pandas DataFrame**. Secondly, The data will be assessed both visually and programmatically. Thirdly, The data will be cleaned, and then the three DataFrames will be merged into one. Finally, I going to anayze and vsiulaize the data.

```
In [1]: #import the required libraries for the EDA  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Data Gathering

```
In [2]: # Convert CSV file into Panadas Dataframe  
benin_df = pd.read_csv("../data/benin-malanville.csv")  
sierraleone_df = pd.read_csv("../data/sierraleone-bumbuna.csv")  
togo_df = pd.read_csv("../data/togo-dapaong_qc.csv")
```

Assessing Data

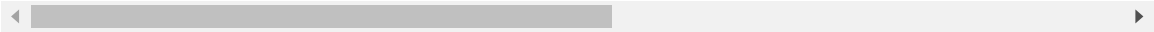
Visual Assessemnt

```
In [3]: benin_df
```

Out[3]:

| | Timestamp | GHI | DNI | DHI | ModA | ModB | Tamb | RH | WS | WSgust | WSstdev |
|--------|------------------|------|------|------|------|------|------|------|-----|--------|---------|
| 0 | 2021-08-09 00:01 | -1.2 | -0.2 | -1.1 | 0.0 | 0.0 | 26.2 | 93.4 | 0.0 | 0.4 | 0.1 |
| 1 | 2021-08-09 00:02 | -1.1 | -0.2 | -1.1 | 0.0 | 0.0 | 26.2 | 93.6 | 0.0 | 0.0 | 0.0 |
| 2 | 2021-08-09 00:03 | -1.1 | -0.2 | -1.1 | 0.0 | 0.0 | 26.2 | 93.7 | 0.3 | 1.1 | 0.5 |
| 3 | 2021-08-09 00:04 | -1.1 | -0.1 | -1.0 | 0.0 | 0.0 | 26.2 | 93.3 | 0.2 | 0.7 | 0.4 |
| 4 | 2021-08-09 00:05 | -1.0 | -0.1 | -1.0 | 0.0 | 0.0 | 26.2 | 93.3 | 0.1 | 0.7 | 0.3 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 525595 | 2022-08-08 23:56 | -5.5 | -0.1 | -5.9 | 0.0 | 0.0 | 23.1 | 98.3 | 0.3 | 1.1 | 0.5 |
| 525596 | 2022-08-08 23:57 | -5.5 | -0.1 | -5.8 | 0.0 | 0.0 | 23.1 | 98.3 | 0.2 | 0.7 | 0.4 |
| 525597 | 2022-08-08 23:58 | -5.5 | -0.1 | -5.8 | 0.0 | 0.0 | 23.1 | 98.4 | 0.6 | 1.1 | 0.5 |
| 525598 | 2022-08-08 23:59 | -5.5 | -0.1 | -5.8 | 0.0 | 0.0 | 23.1 | 98.3 | 0.9 | 1.3 | 0.5 |
| 525599 | 2022-08-09 00:00 | -5.5 | -0.1 | -5.7 | 0.0 | 0.0 | 23.1 | 98.3 | 1.2 | 1.6 | 0.5 |

525600 rows × 19 columns



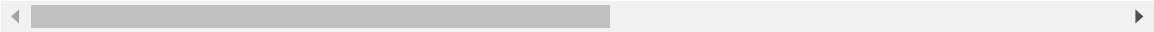
- Missing values in the `comments` column
- `GHI` , `DNI` , and `DHI` columns contains negative values

In [4]: `sierraleone_df`

Out[4]:

| | Timestamp | GHI | DNI | DHI | ModA | ModB | Tamb | RH | WS | WSgust | WSstdc |
|--------|------------------|------|------|------|------|------|------|-------|-----|--------|--------|
| 0 | 2021-10-30 00:01 | -0.7 | -0.1 | -0.8 | 0.0 | 0.0 | 21.9 | 99.1 | 0.0 | 0.0 | 0 |
| 1 | 2021-10-30 00:02 | -0.7 | -0.1 | -0.8 | 0.0 | 0.0 | 21.9 | 99.2 | 0.0 | 0.0 | 0 |
| 2 | 2021-10-30 00:03 | -0.7 | -0.1 | -0.8 | 0.0 | 0.0 | 21.9 | 99.2 | 0.0 | 0.0 | 0 |
| 3 | 2021-10-30 00:04 | -0.7 | 0.0 | -0.8 | 0.0 | 0.0 | 21.9 | 99.3 | 0.0 | 0.0 | 0 |
| 4 | 2021-10-30 00:05 | -0.7 | -0.1 | -0.8 | 0.0 | 0.0 | 21.9 | 99.3 | 0.0 | 0.0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 525595 | 2022-10-29 23:56 | -1.6 | -0.1 | -2.9 | 0.0 | 0.0 | 24.0 | 100.0 | 0.0 | 0.0 | 0 |
| 525596 | 2022-10-29 23:57 | -1.7 | -0.1 | -3.0 | 0.0 | 0.0 | 24.0 | 100.0 | 0.0 | 0.0 | 0 |
| 525597 | 2022-10-29 23:58 | -1.7 | -0.1 | -3.1 | 0.0 | 0.0 | 24.0 | 100.0 | 0.0 | 0.0 | 0 |
| 525598 | 2022-10-29 23:59 | -1.7 | -0.2 | -3.3 | 0.0 | 0.0 | 23.9 | 100.0 | 0.0 | 0.0 | 0 |
| 525599 | 2022-10-30 00:00 | -1.7 | -0.1 | -3.4 | 0.0 | 0.0 | 23.9 | 100.0 | 0.0 | 0.0 | 0 |

525600 rows × 19 columns



- Missing values in the `comments` column
- `GHI` , `DNI` , and `DHI` columns contains negative values

In [5]:

```
togo_df
```

Out[5]:

| | Timestamp | GHI | DNI | DHI | ModA | ModB | Tamb | RH | WS | WSgust | WSstdev |
|--------|------------------|------|-----|-----|------|------|------|------|-----|--------|---------|
| 0 | 2021-10-25 00:01 | -1.3 | 0.0 | 0.0 | 0.0 | 0.0 | 24.8 | 94.5 | 0.9 | 1.1 | 0.4 |
| 1 | 2021-10-25 00:02 | -1.3 | 0.0 | 0.0 | 0.0 | 0.0 | 24.8 | 94.4 | 1.1 | 1.6 | 0.4 |
| 2 | 2021-10-25 00:03 | -1.3 | 0.0 | 0.0 | 0.0 | 0.0 | 24.8 | 94.4 | 1.2 | 1.4 | 0.3 |
| 3 | 2021-10-25 00:04 | -1.2 | 0.0 | 0.0 | 0.0 | 0.0 | 24.8 | 94.3 | 1.2 | 1.6 | 0.3 |
| 4 | 2021-10-25 00:05 | -1.2 | 0.0 | 0.0 | 0.0 | 0.0 | 24.8 | 94.0 | 1.3 | 1.6 | 0.4 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 525595 | 2022-10-24 23:56 | -0.8 | 0.0 | 0.0 | 0.0 | 0.0 | 25.2 | 53.8 | 0.0 | 0.0 | 0.0 |
| 525596 | 2022-10-24 23:57 | -0.9 | 0.0 | 0.0 | 0.0 | 0.0 | 25.3 | 53.5 | 0.0 | 0.0 | 0.0 |
| 525597 | 2022-10-24 23:58 | -1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 25.3 | 53.4 | 0.0 | 0.0 | 0.0 |
| 525598 | 2022-10-24 23:59 | -1.1 | 0.0 | 0.0 | 0.0 | 0.0 | 25.4 | 53.5 | 0.0 | 0.0 | 0.0 |
| 525599 | 2022-10-25 00:00 | -1.2 | 0.0 | 0.0 | 0.0 | 0.0 | 25.4 | 52.3 | 0.0 | 0.0 | 0.0 |

525600 rows × 19 columns



- Missing values in the `comments` column
- `GHI` column contains negative values

Programmatic Assessement

```
In [6]: benin_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525600 entries, 0 to 525599
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Timestamp              525600 non-null object
1   GHI                    525600 non-null float64
2   DNI                    525600 non-null float64
3   DHI                    525600 non-null float64
4   ModA                   525600 non-null float64
5   ModB                   525600 non-null float64
6   Tamb                   525600 non-null float64
7   RH                     525600 non-null float64
8   WS                     525600 non-null float64
9   WSgust                 525600 non-null float64
10  WSstdev                525600 non-null float64
11  WD                     525600 non-null float64
12  WDstdev                525600 non-null float64
13  BP                     525600 non-null int64
14  Cleaning               525600 non-null int64
15  Precipitation          525600 non-null float64
16  TModA                  525600 non-null float64
17  TModB                  525600 non-null float64
18  Comments                0 non-null      float64
dtypes: float64(16), int64(2), object(1)
memory usage: 76.2+ MB
```

- Erroneous datatype `Timestamp` column should be `Datetime`

```
In [7]: #check duplicated values in benin_df
sum(benin_df.duplicated())
```

```
Out[7]: 0
```

```
In [8]: # check null values in benin_df
benin_df.isna().sum()
```

```
Out[8]: Timestamp              0
GHI                          0
DNI                          0
DHI                          0
ModA                         0
ModB                         0
Tamb                         0
RH                           0
WS                           0
WSgust                       0
WSstdev                      0
WD                           0
WDstdev                      0
BP                           0
Cleaning                     0
Precipitation                0
TModA                        0
TModB                        0
Comments                     525600
dtype: int64
```

```
In [9]: # Summary Statistics
        benin_df.describe()
```

```
Out[9]:
```

| | GHI | DNI | DHI | ModA | ModB |
|-------|---------------|---------------|---------------|---------------|---------------|
| count | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 |
| mean | 240.559452 | 167.187516 | 115.358961 | 236.589496 | 228.883576 |
| std | 331.131327 | 261.710501 | 158.691074 | 326.894859 | 316.536515 |
| min | -12.900000 | -7.800000 | -12.600000 | 0.000000 | 0.000000 |
| 25% | -2.000000 | -0.500000 | -2.100000 | 0.000000 | 0.000000 |
| 50% | 1.800000 | -0.100000 | 1.600000 | 4.500000 | 4.300000 |
| 75% | 483.400000 | 314.200000 | 216.300000 | 463.700000 | 447.900000 |
| max | 1413.000000 | 952.300000 | 759.200000 | 1342.300000 | 1342.300000 |

```
In [10]: sierraleone_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525600 entries, 0 to 525599
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Timestamp              525600 non-null object
1   GHI                    525600 non-null float64
2   DNI                    525600 non-null float64
3   DHI                    525600 non-null float64
4   ModA                   525600 non-null float64
5   ModB                   525600 non-null float64
6   Tamb                   525600 non-null float64
7   RH                     525600 non-null float64
8   WS                     525600 non-null float64
9   WSGust                 525600 non-null float64
10  WStddev                 525600 non-null float64
11  WD                     525600 non-null float64
12  WDstddev                525600 non-null float64
13  BP                     525600 non-null int64
14  Cleaning                525600 non-null int64
15  Precipitation           525600 non-null float64
16  TModA                   525600 non-null float64
17  TModB                   525600 non-null float64
18  Comments                 0 non-null      float64
dtypes: float64(16), int64(2), object(1)
memory usage: 76.2+ MB
```

- Erroneous datatype `Timestamp` column should be `Datetime`

```
In [11]: #check duplicated values in benin_df
        sum(sierraleone_df.duplicated())
```

```
Out[11]: 0
```

```
In [12]: # check null values in benin_df
        sierraleone_df.isna().sum()
```

```
Out[12]: Timestamp      0
        GHI             0
        DNI             0
        DHI             0
        ModA            0
        ModB            0
        Tamb            0
        RH              0
        WS              0
        WSgust          0
        WSstdev         0
        WD              0
        WDstdev         0
        BP              0
        Cleaning        0
        Precipitation    0
        TModA           0
        TModB           0
        Comments        525600
        dtype: int64
```

```
In [13]: # Summary Statistics
        sierraleone_df.describe()
```

Out[13]:

| | GHI | DNI | DHI | ModA | ModB |
|-------|---------------|---------------|---------------|---------------|---------------|
| count | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 |
| mean | 201.957515 | 116.376337 | 113.720571 | 206.643095 | 198.114691 |
| std | 298.495150 | 218.652659 | 158.946032 | 300.896893 | 288.889073 |
| min | -19.500000 | -7.800000 | -17.900000 | 0.000000 | 0.000000 |
| 25% | -2.800000 | -0.300000 | -3.800000 | 0.000000 | 0.000000 |
| 50% | 0.300000 | -0.100000 | -0.100000 | 3.600000 | 3.400000 |
| 75% | 362.400000 | 107.000000 | 224.700000 | 359.500000 | 345.400000 |
| max | 1499.000000 | 946.000000 | 892.000000 | 1507.000000 | 1473.000000 |

```
In [14]: togo_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525600 entries, 0 to 525599
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Timestamp              525600 non-null object
1   GHI                    525600 non-null float64
2   DNI                    525600 non-null float64
3   DHI                    525600 non-null float64
4   ModA                   525600 non-null float64
5   ModB                   525600 non-null float64
6   Tamb                   525600 non-null float64
7   RH                     525600 non-null float64
8   WS                     525600 non-null float64
9   WSgust                 525600 non-null float64
10  WSstdev                525600 non-null float64
11  WD                     525600 non-null float64
12  WDstdev                525600 non-null float64
13  BP                     525600 non-null int64
14  Cleaning               525600 non-null int64
15  Precipitation          525600 non-null float64
16  TModA                  525600 non-null float64
17  TModB                  525600 non-null float64
18  Comments                0 non-null      float64
dtypes: float64(16), int64(2), object(1)
memory usage: 76.2+ MB
```

- Erroneous datatype `Timestamp` column should be `Datetime`

```
In [15]: #check duplicated values in benin_df
sum(togo_df.duplicated())
```

```
Out[15]: 0
```

```
In [16]: # check null values in benin_df
togo_df.isna().sum()
```

```
Out[16]: Timestamp          0
GHI                        0
DNI                        0
DHI                        0
ModA                       0
ModB                       0
Tamb                       0
RH                          0
WS                          0
WSgust                     0
WSstdev                    0
WD                          0
WDstdev                    0
BP                          0
Cleaning                   0
Precipitation              0
TModA                      0
TModB                      0
Comments                   525600
dtype: int64
```



```
In [17]: # Summary Statistics
togo_df.describe()
```

Out[17]:

| | GHI | DNI | DHI | ModA | ModB |
|-------|---------------|---------------|---------------|---------------|---------------|
| count | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 | 525600.000000 |
| mean | 230.555040 | 151.258469 | 116.444352 | 226.144375 | 219.568588 |
| std | 322.532347 | 250.956962 | 156.520714 | 317.346938 | 307.932510 |
| min | -12.700000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | -2.200000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 50% | 2.100000 | 0.000000 | 2.500000 | 4.400000 | 4.300000 |
| 75% | 442.400000 | 246.400000 | 215.700000 | 422.525000 | 411.000000 |
| max | 1424.000000 | 1004.500000 | 805.700000 | 1380.000000 | 1367.000000 |

Quality issues

1. Missing values in the `comments` column.
2. `GHI`, `DNI`, and `DHI` columns contains negative values.
3. Erroneous datatype `Timestamp` column should be `Datetime`.

Data Cleaning

Quality issues

```
In [18]: # make a copies of the original data
benin_df_clean = benin_df.copy()
sierraleone_df_clean = sierraleone_df.copy()
togo_df_clean = togo_df.copy()
```

Issue #1:

The `comments` column contains missing values

Define

The `comments` column contains entirely null value in all the three dataframe. I use pandas `drop` method to drop the column.

Code

```
In [19]: benin_df_clean.drop("Comments", axis = 1, inplace=True)
sierraleone_df_clean.drop("Comments", axis = 1, inplace=True)
togo_df_clean.drop("Comments", axis = 1, inplace=True)
```

Test

```
In [20]: print(benin_df_clean.columns)
print(sierraleone_df_clean.columns)
print(togo_df_clean.columns)

Index(['Timestamp', 'GHI', 'DNI', 'DHI', 'ModA', 'ModB', 'Tamb', 'RH', 'WS',
      'WSgust', 'WSstdev', 'WD', 'WDstdev', 'BP', 'Cleaning', 'Precipitation',
      'TModA', 'TModB'],
      dtype='object')
Index(['Timestamp', 'GHI', 'DNI', 'DHI', 'ModA', 'ModB', 'Tamb', 'RH', 'WS',
      'WSgust', 'WSstdev', 'WD', 'WDstdev', 'BP', 'Cleaning', 'Precipitation',
      'TModA', 'TModB'],
      dtype='object')
Index(['Timestamp', 'GHI', 'DNI', 'DHI', 'ModA', 'ModB', 'Tamb', 'RH', 'WS',
      'WSgust', 'WSstdev', 'WD', 'WDstdev', 'BP', 'Cleaning', 'Precipitation',
      'TModA', 'TModB'],
      dtype='object')
```

Issue #2:

GHI, **DNI**, and **DHI** columns contains negative values.

Define

GHI, **DNI**, and **DHI** columns contains negative values in all the three dataframe. I use 'abs' function to convert each value into absolute value.

Code

```
In [21]: benin_df_clean[['GHI', 'DNI', 'DHI']] = benin_df_clean[['GHI', 'DNI', 'DHI']].abs()
sierraleone_df_clean[['GHI', 'DNI', 'DHI']] = sierraleone_df_clean[['GHI', 'DNI', 'DHI']].abs()
togo_df_clean[['GHI', 'DNI', 'DHI']] = togo_df_clean[['GHI', 'DNI', 'DHI']].abs()
```

Test

```
In [22]: benin_df_clean[['GHI', 'DNI', 'DHI']]
```

Out[22]:

| | GHI | DNI | DHI |
|--------|-----|-----|-----|
| 0 | 1.2 | 0.2 | 1.1 |
| 1 | 1.1 | 0.2 | 1.1 |
| 2 | 1.1 | 0.2 | 1.1 |
| 3 | 1.1 | 0.1 | 1.0 |
| 4 | 1.0 | 0.1 | 1.0 |
| ... | ... | ... | ... |
| 525595 | 5.5 | 0.1 | 5.9 |
| 525596 | 5.5 | 0.1 | 5.8 |
| 525597 | 5.5 | 0.1 | 5.8 |
| 525598 | 5.5 | 0.1 | 5.8 |
| 525599 | 5.5 | 0.1 | 5.7 |

525600 rows × 3 columns

```
In [23]: sierraleone_df_clean[['GHI', 'DNI', 'DHI']]
```

Out[23]:

| | GHI | DNI | DHI |
|--------|-----|-----|-----|
| 0 | 0.7 | 0.1 | 0.8 |
| 1 | 0.7 | 0.1 | 0.8 |
| 2 | 0.7 | 0.1 | 0.8 |
| 3 | 0.7 | 0.0 | 0.8 |
| 4 | 0.7 | 0.1 | 0.8 |
| ... | ... | ... | ... |
| 525595 | 1.6 | 0.1 | 2.9 |
| 525596 | 1.7 | 0.1 | 3.0 |
| 525597 | 1.7 | 0.1 | 3.1 |
| 525598 | 1.7 | 0.2 | 3.3 |
| 525599 | 1.7 | 0.1 | 3.4 |

525600 rows × 3 columns

```
In [24]: togo_df_clean[['GHI', 'DNI', 'DHI']]
```

Out[24]:

| | GHI | DNI | DHI |
|--------|-----|-----|-----|
| 0 | 1.3 | 0.0 | 0.0 |
| 1 | 1.3 | 0.0 | 0.0 |
| 2 | 1.3 | 0.0 | 0.0 |
| 3 | 1.2 | 0.0 | 0.0 |
| 4 | 1.2 | 0.0 | 0.0 |
| ... | ... | ... | ... |
| 525595 | 0.8 | 0.0 | 0.0 |
| 525596 | 0.9 | 0.0 | 0.0 |
| 525597 | 1.0 | 0.0 | 0.0 |
| 525598 | 1.1 | 0.0 | 0.0 |
| 525599 | 1.2 | 0.0 | 0.0 |

525600 rows × 3 columns

Issue #3:

Erroneous datatype `Timestamp` column should be Datetime.

Define

`Timestamp` column contains wrong datatype in all the three dataframes, it should has a datetime datatype. I use `to_datetime` pandas method to convert `object` type into `datetime` type.

Code

```
In [26]: # convert the datatype
benin_df_clean['Timestamp'] = pd.to_datetime(benin_df_clean['Timestamp'])
sierraleone_df_clean['Timestamp'] = pd.to_datetime(sierraleone_df_clean['Timestamp'])
togo_df_clean['Timestamp'] = pd.to_datetime(togo_df_clean['Timestamp']).d
```

Test

```
In [27]: # check the datatype
print(benin_df_clean['Timestamp'].dtype)
print(sierraleone_df_clean['Timestamp'].dtype)
print(togo_df_clean['Timestamp'].dtype)
```

```
datetime64[ns]
datetime64[ns]
datetime64[ns]
```

Concatination

Concatinate all the three dataframe into one dataframe.

```
In [28]: all_clean_df = pd.concat([benin_df_clean, sierraleone_df_clean, togo_df_clean], axis=0)
all_clean_df.sample(5)
```

```
Out[28]:
```

| | Timestamp | GHI | DNI | DHI | ModA | ModB | Tamb | RH | WS | WSgust | W |
|--------|------------|-------|-------|-------|-------|-------|------|-------|-----|--------|---|
| 305798 | 2022-05-30 | 242.4 | 27.4 | 228.5 | 226.9 | 221.6 | 24.4 | 93.7 | 0.0 | 0.0 | |
| 325883 | 2022-06-13 | 162.4 | 359.0 | 86.4 | 109.7 | 104.6 | 23.4 | 99.3 | 0.0 | 0.0 | |
| 117825 | 2022-01-19 | 13.0 | 0.6 | 13.0 | 0.0 | 0.0 | 24.6 | 51.0 | 0.0 | 0.0 | |
| 479717 | 2022-09-28 | 0.8 | 0.1 | 3.7 | 0.0 | 0.0 | 22.2 | 99.7 | 0.0 | 0.0 | |
| 24908 | 2021-11-16 | 30.4 | 0.1 | 30.3 | 31.9 | 30.2 | 23.3 | 100.0 | 0.0 | 0.0 | |

```
In [33]: all_clean_df.shape
```

```
Out[33]: (1576800, 18)
```

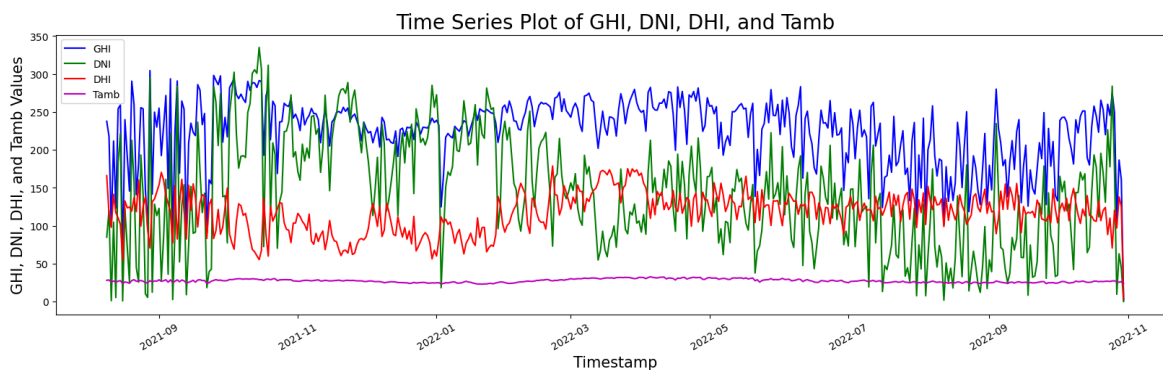
Time Series Analysis

Analyze how variables like GHI, DNI, DHI, and Tamb change over time.

```
In [29]: # group data by Timestamp column for Timeseries plot
grouped_df = all_clean_df.groupby(['Timestamp']).mean()[['GHI', 'DNI', 'DHI', 'Tamb']]
grouped_df.reset_index(inplace=True)
```

```
In [30]: cols = ['GHI', 'DNI', 'DHI', 'Tamb']
colors = ['b', 'g', 'r', 'm'] # blue, green, red, magenta
fig = plt.subplots(figsize=(20, 5))
for col, color in zip(cols, colors):
    plt.plot(grouped_df['Timestamp'], grouped_df[col], label=col, color=color)
plt.title('Time Series Plot of GHI, DNI, DHI, and Tamb', fontsize=20)
plt.xlabel('Timestamp', fontsize=15)
plt.ylabel('GHI, DNI, DHI, and Tamb Values', fontsize=15)
plt.xticks(rotation=30)
plt.legend()
```

```
Out[30]: <matplotlib.legend.Legend at 0x74d3d246e950>
```



As we can see from the above graph, the **Tamb** data seems fixed throughout all the years, and the rest of the column data (**DHI** , **GHI** , **DNI**) vary throughout the years.

Correlation Analysis

Determine the correlation between different variables like solar radiation components (**GHI**, **DHI**, **DNI**) and temperature measures (**TModA**, **TModB**) to uncover relationships.

To analyze the relationships between variables, I employ correlation coefficients calculated using the **corr** method of pandas DataFrames.

```
In [31]: #Calculate correlation coefficients
correlation_matrix = all_clean_df[["GHI", "DHI", "DNI", "TModA", "TModB"]]
print(correlation_matrix)
```

| | GHI | DHI | DNI | TModA | TModB |
|-------|----------|----------|----------|----------|----------|
| GHI | 1.000000 | 0.851051 | 0.876611 | 0.905345 | 0.898338 |
| DHI | 0.851051 | 1.000000 | 0.531800 | 0.800149 | 0.797958 |
| DNI | 0.876611 | 0.531800 | 1.000000 | 0.784409 | 0.775983 |
| TModA | 0.905345 | 0.800149 | 0.784409 | 1.000000 | 0.969891 |
| TModB | 0.898338 | 0.797958 | 0.775983 | 0.969891 | 1.000000 |

As we can see from the correlation coefficient results, there is a strong relationship between the solar radiation components (**GHI**, **DHI**, **DNI**) and temperature measures (**TModA**, **TModB**). Notably, **GHI** shows the strongest correlation with the temperature measures (**TModA** = 0.905345 and **TModB** = 0.898338).

Wind Analysis

Explore wind speed (**WS**, **WSgust**, **WSstdev**) and wind direction (**WD**, **WDstdev**) data to identify any trends or notable wind events.

I employ summary statistics (**mean**, **median**, **standard deviation**, **minimum**, **maximum**) for wind speed (**WS**, **WSgust**, and **WSstdev**) to understand central tendency and variability.

```
In [33]: wind_speed_stats = all_clean_df[["WS", "WSgust", "WSstdev"]].describe()
print(wind_speed_stats)
```

| | WS | WSgust | WSstdev |
|-------|--------------|--------------|--------------|
| count | 1.576800e+06 | 1.576800e+06 | 1.576800e+06 |
| mean | 1.878440e+00 | 2.576763e+00 | 4.649840e-01 |
| std | 1.536357e+00 | 1.961275e+00 | 2.904002e-01 |
| min | 0.000000e+00 | 0.000000e+00 | 0.000000e+00 |
| 25% | 6.000000e-01 | 1.100000e+00 | 4.000000e-01 |
| 50% | 1.800000e+00 | 2.400000e+00 | 5.000000e-01 |
| 75% | 2.800000e+00 | 3.600000e+00 | 6.000000e-01 |
| max | 1.950000e+01 | 2.660000e+01 | 4.700000e+00 |

Based on the summary statistics output:

- The average wind speed is **1.88 m/s**, indicating a moderate wind regime.
- The standard deviation of wind speed is **1.54 m/s**, suggesting a moderate variability in wind speeds around the mean.
- The wind speed ranged from **0 m/s** (calm winds) to a maximum of **19.5 m/s** (strong breeze).

- The quartiles (25th and 75th percentiles) show that wind speeds are distributed between 0.6 m/s and 2.8 m/s for most of the time (IQR). The 50th percentile (median) is 1.8 m/s, which is close to the mean, suggesting a symmetrical distribution.

```
In [40]: wind_data = all_clean_df.groupby(['Timestamp']).mean()[['WS', 'WSgust', 'WD', 'WDstddev']]
wind_data.reset_index(inplace=True)
```

```
In [41]: # Plotting time-series graphs for each column
plt.figure(figsize=(14, 10))

# WS
plt.subplot(3, 2, 1)
plt.plot(wind_data['Timestamp'], wind_data['WS'], color='blue')
plt.title('Wind Speed (WS)')
plt.xlabel('Time')
plt.ylabel('Wind Speed (m/s)')

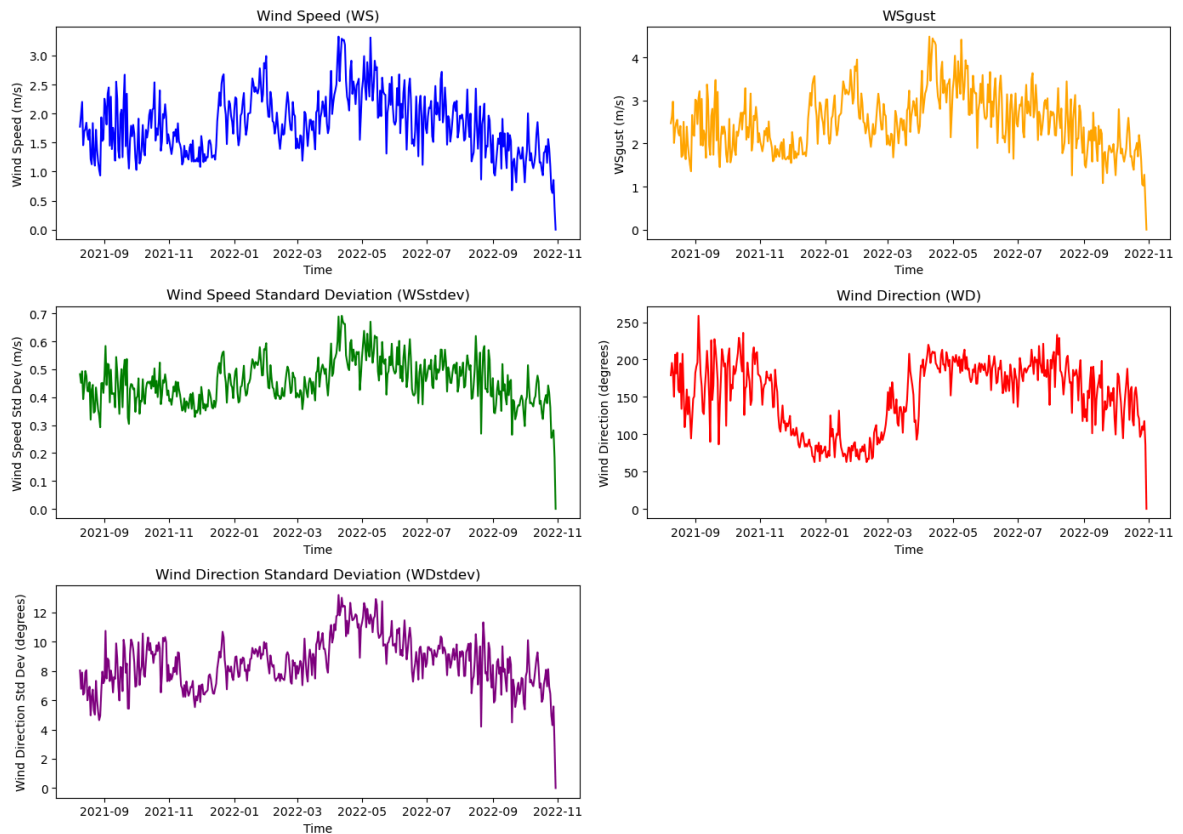
# WSgust
plt.subplot(3, 2, 2)
plt.plot(wind_data['Timestamp'], wind_data['WSgust'], color='orange')
plt.title('WSgust')
plt.xlabel('Time')
plt.ylabel('WSgust (m/s)')

# WSstddev
plt.subplot(3, 2, 3)
plt.plot(wind_data['Timestamp'], wind_data['WSstddev'], color='green')
plt.title('Wind Speed Standard Deviation (WSstddev)')
plt.xlabel('Time')
plt.ylabel('Wind Speed Std Dev (m/s)')

# WD
plt.subplot(3, 2, 4)
plt.plot(wind_data['Timestamp'], wind_data['WD'], color='red')
plt.title('Wind Direction (WD)')
plt.xlabel('Time')
plt.ylabel('Wind Direction (degrees)')

# WDstddev
plt.subplot(3, 2, 5)
plt.plot(wind_data['Timestamp'], wind_data['WDstddev'], color='purple')
plt.title('Wind Direction Standard Deviation (WDstddev)')
plt.xlabel('Time')
plt.ylabel('Wind Direction Std Dev (degrees)')

plt.tight_layout()
```



The time-series graphs reveal that wind speed and wind direction have the same pattern.

Temperature Analysis

Compare module temperatures (`TModA` , `TModB`) with ambient temperature (`Tamb`) to see how they are related or vary under different conditions.

I calculate the correlation coefficient between each module temperature and ambient temperature. This will quantify the strength and direction of the linear relationship.

```
In [43]: # calculate the correlation coefficient
correlation_matrix = all_clean_df[["TModA", "TModB", "Tamb"]].corr()
print(correlation_matrix)
```

| | TModA | TModB | Tamb |
|-------|----------|----------|----------|
| TModA | 1.000000 | 0.969891 | 0.787788 |
| TModB | 0.969891 | 1.000000 | 0.789931 |
| Tamb | 0.787788 | 0.789931 | 1.000000 |

```
In [46]: # plotting Scatter plot

plt.figure(figsize=(14, 10))

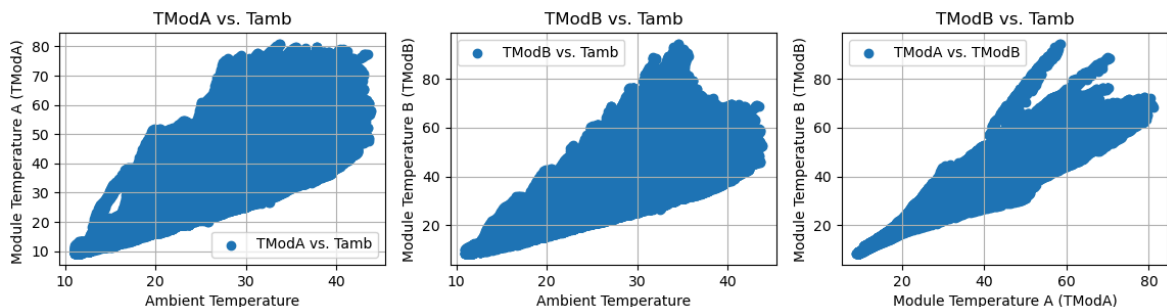
#TModA vs Tamb
plt.subplot(3, 3, 1)
plt.scatter(all_clean_df["Tamb"], all_clean_df["TModA"], label="TModA vs. Tamb")
plt.xlabel("Ambient Temperature")
plt.ylabel("Module Temperature A (TModA) ")
plt.title("TModA vs. Tamb")
plt.grid(True)
plt.legend()
```



```
# TModB vs Tamb
plt.subplot(3, 3, 2)
plt.scatter(all_clean_df["Tamb"], all_clean_df["TModB"], label="TModB vs. Tamb")
plt.xlabel("Ambient Temperature")
plt.ylabel("Module Temperature B (TModB)")
plt.title("TModB vs. Tamb")
plt.grid(True)
plt.legend()

# TModA vs TModB
plt.subplot(3, 3, 3)
plt.scatter(all_clean_df["TModA"], all_clean_df["TModB"], label="TModA vs. TModB")
plt.xlabel("Module Temperature A (TModA)")
plt.ylabel("Module Temperature B (TModB)")
plt.title("TModA vs. TModB")
plt.grid(True)
plt.legend()
```

Out[46]: <matplotlib.legend.Legend at 0x74d395b0b190>



Both the correlation coefficient and the scatter plot output shows there relationships between the module temperatures (TModA , TModB) and ambient temperature (Tamb).

- The correlation coefficient of **0.97** indicates a very strong positive linear relationship between TModA and TModB . This suggests that both module temperatures tend to move in the same direction and experience similar changes.
- The correlation coefficients between module temperatures (around **0.79**) and ambient temperature (Tamb) are positive, indicating a tendency for module temperatures to increase as ambient temperature increases. However, the value is not as high as the correlation between the modules themselves.

Histograms

Create histograms for variables like GHI , DNI , DHI , WS , and temperatures to visualize the frequency distribution of these variables.

```
In [47]: # Histograms plot
plt.figure(figsize=(14, 10))

plt.subplot(2, 3, 1)
sns.histplot(all_clean_df['GHI'], bins=20, kde=True, color='blue')
plt.title('Global Horizontal Irradiance (GHI) Histogram')

plt.subplot(2, 3, 2)
```

```
sns.histplot(all_clean_df['DNI'], bins=20, kde=True, color='orange')
plt.title('Direct Normal Irradiance (DNI) Histogram')

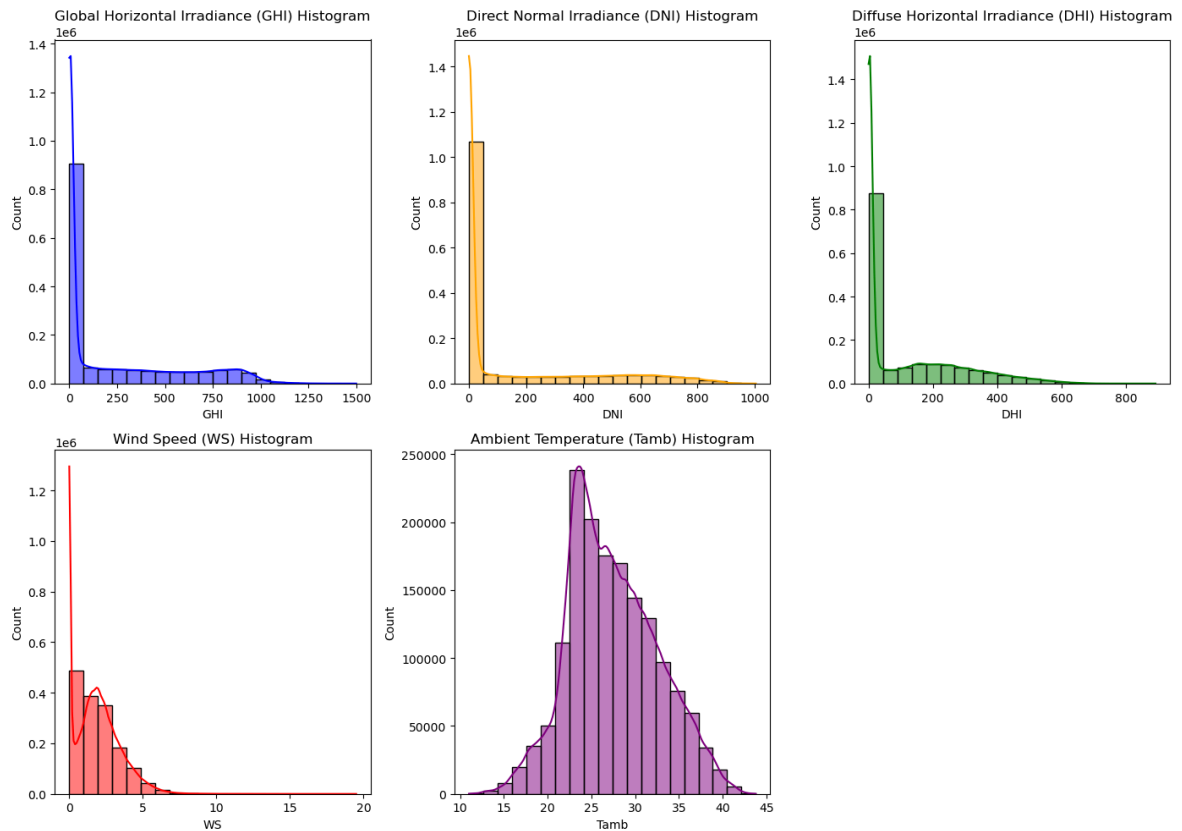
plt.subplot(2, 3, 3)
sns.histplot(all_clean_df['DHI'], bins=20, kde=True, color='green')
plt.title('Diffuse Horizontal Irradiance (DHI) Histogram')

plt.subplot(2, 3, 4)
sns.histplot(all_clean_df['WS'], bins=20, kde=True, color='red')
plt.title('Wind Speed (WS) Histogram')

plt.subplot(2, 3, 5)
sns.histplot(all_clean_df['Tamb'], bins=20, kde=True, color='purple')
plt.title('Ambient Temperature (Tamb) Histogram')

plt.tight_layout()
plt.show()
```

```
/home/derbew/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:11
19: FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/home/derbew/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:11
19: FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/home/derbew/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:11
19: FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/home/derbew/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:11
19: FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/home/derbew/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:11
19: FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```



As seen from the above histogram plot, the parameters **GHI**, **DNI**, **DHI**, and **WS** exhibit right-skewed (positive skewness) distributions. In contrast, the **Tamb** parameter has a symmetrical distribution, suggesting that its data is evenly distributed on both sides of the central tendency.

Box Plots

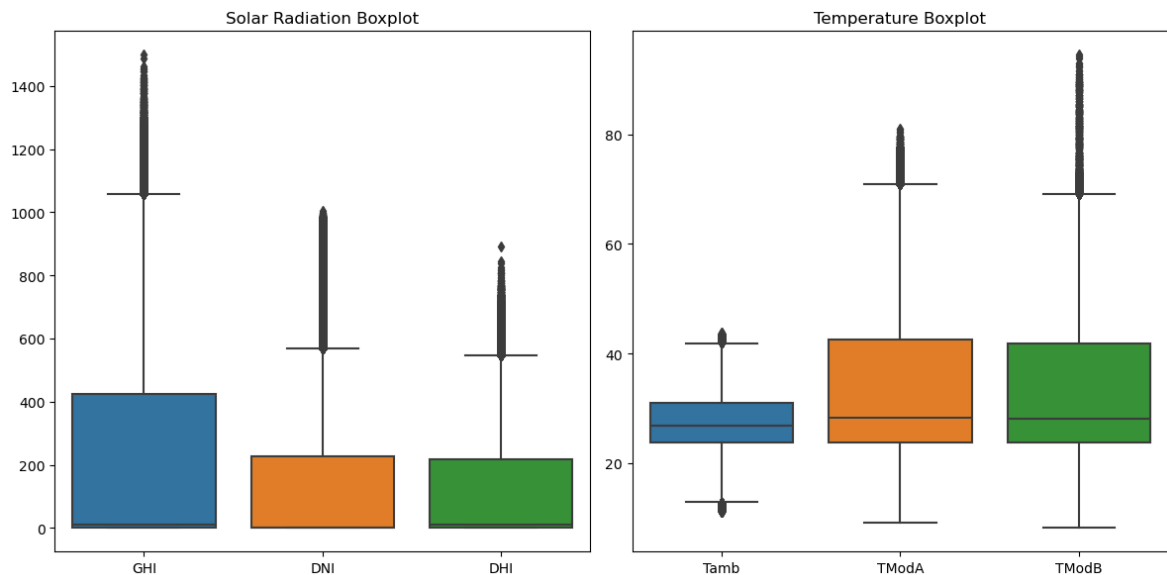
Use box plots to examine the spread and presence of outliers in the solar radiation and temperature data.

```
In [48]: # Box Plots
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
sns.boxplot(data=all_clean_df[['GHI', 'DNI', 'DHI']])
plt.title('Solar Radiation Boxplot')

plt.subplot(1, 2, 2)
sns.boxplot(data=all_clean_df[['Tamb', 'TModA', 'TModB']])
plt.title('Temperature Boxplot')

plt.tight_layout()
plt.show()
```



Scatter Plots

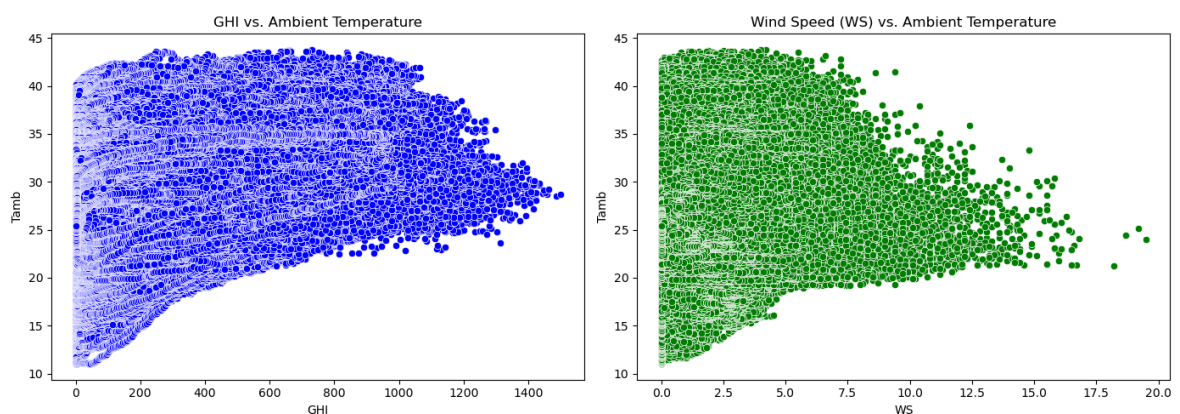
Generate scatter plots to explore the relationships between pairs of variables, such as **GHI** vs. **Tamb**, **WS** vs. **WSgust**, or any other potentially interesting pairs.

```
In [49]: # Scatter Plots
plt.figure(figsize=(14, 5))

plt.subplot(1, 2, 1)
sns.scatterplot(x='GHI', y='Tamb', data=all_clean_df, color='blue')
plt.title('GHI vs. Ambient Temperature')

plt.subplot(1, 2, 2)
sns.scatterplot(x='WS', y='Tamb', data=all_clean_df, color='green')
plt.title('Wind Speed (WS) vs. Ambient Temperature')

plt.tight_layout()
plt.show()
```



As we can see from the above scatter plot, there appears to be no correlation between **GHI** and **Tamb**, nor between **WS** and **Tamb**.

Conclusions

This project aimed to practice exploratory data analysis (EDA) skills using real-world solar radiation measurement data collected from three African countries: Benin, Sierra Leone, and Togo. The EDA process consisted of three phases:

1. Data Gathering: The data was provided by 10 Academy.
2. Data Cleaning and Preprocessing: The data was cleaned and prepared for analysis.
3. Data Analysis and Visualization: The data was analyzed and visualized to extract insights.

Here is the Key Insights:

- A strong positive correlation exists between module temperatures (TModA and TModB).
- A moderate positive correlation exists between module and ambient temperatures.
- There is no correlation between GHI and Tamb, nor between WS and Tamb.
- The parameters GHI, DNI, DHI, and WS exhibit right-skewed (positive skewness) distributions.
- The Tamb parameter has a symmetrical distribution.
- The time-series graphs reveal that wind speed and wind direction exhibit a similar pattern.

Limitations

Due to time limitations, I was only able to assess three key aspects of the dataset. A more comprehensive analysis would require additional time to explore all potential issues. In the analysis and visualization stage, I focused on obtaining a few initial insights from the dataset. Further exploration through EDA can reveal additional valuable information.

References

[Pandas Documentation](#)

[Matplotlib Documentation](#)

[Seaborn Documentation](#)

[Exploratory Data Analysis](#)

In []: